

Insurance FAQ and Knowledge Governance

Synthetic educational content created for an insurance knowledge-copilot prototype. Verify all real-world decisions against official policy wording, insurer procedures, and applicable regulations.

Purpose

This synthetic FAQ supports retrieval and governance testing. It is not legal, regulatory, underwriting, or claims advice.

What is a policy schedule?

It summarizes key details such as the insured party, policy period, coverage, limits, deductibles, and premium. It should be read with the policy wording and endorsements.

Coverage versus exclusion

Coverage describes events or losses the policy may respond to, subject to conditions. An exclusion identifies losses or circumstances that are restricted or not covered.

Why versions matter

Current versions reduce the risk of relying on outdated terms, procedures, or forms. Documents should have an owner, effective date, review date, and version identifier.

Handling outdated knowledge

Mark outdated content, prevent it from being presented as current where possible, retain it only for justified historical or audit needs, and route it to the owner.

Missing evidence

An AI assistant should say when available sources do not establish an answer, show supporting sources where available, and direct the user to an official document or subject-matter expert.

Metadata

Metadata supports filtering, classification, ownership, lifecycle management, reporting-period selection, and retrieval refinement. Useful fields include type, department, region, product, dates, owner, sensitivity, and source.

Access management

Authorization should be evaluated using trusted identity and access information before content is sent to a language model. The model is not the security boundary.

Governance checklist

Each managed document should have an owner, source, version, effective date, review date, sensitivity classification, retention decision, and escalation path.